

In [1]:



```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
df=pd.read_csv(r"C:\Users\G S R KARTHIK\Downloads\fiat500_VehicleSelection_Dataset.csv")
print(df)
```

	ID	model	engine_power	age_in_days	km	previous_owners
0	1	lounge	51	882	25000	1 \
1	2	pop	51	1186	32500	1
2	3	sport	74	4658	142228	1
3	4	lounge	51	2739	160000	1
4	5	pop	73	3074	106880	1
...
1533	1534	sport	51	3712	115280	1
1534	1535	lounge	74	3835	112000	1
1535	1536	pop	51	2223	60457	1
1536	1537	lounge	51	2557	80750	1
1537	1538	pop	51	1766	54276	1

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]

In [2]:

```
df.head(10)
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
5	6	pop	74	3623	70225	1	45.000702	7.682270
6	7	lounge	51	731	11600	1	44.907242	8.611560
7	8	lounge	51	1521	49076	1	41.903221	12.495650
8	9	sport	73	4049	76000	1	45.548000	11.549470
9	10	sport	51	3653	89000	1	45.438301	10.991700

In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   ID                   1538 non-null   int64
1   model                1538 non-null   object
2   engine_power         1538 non-null   int64
3   age_in_days          1538 non-null   int64
4   km                   1538 non-null   int64
5   previous_owners      1538 non-null   int64
6   lat                  1538 non-null   float64
7   lon                  1538 non-null   float64
8   price                1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [4]:

```
df.describe()
```

Out[4]:

	ID	engine_power	age_in_days	km	previous_owners	li
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.54136
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.13351
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.85583
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.80295
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.39405
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.46796
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.79561

In [5]:

```
df.columns
```

Out[5]:

```
Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owner  
s',  
      'lat', 'lon', 'price'],  
      dtype='object')
```

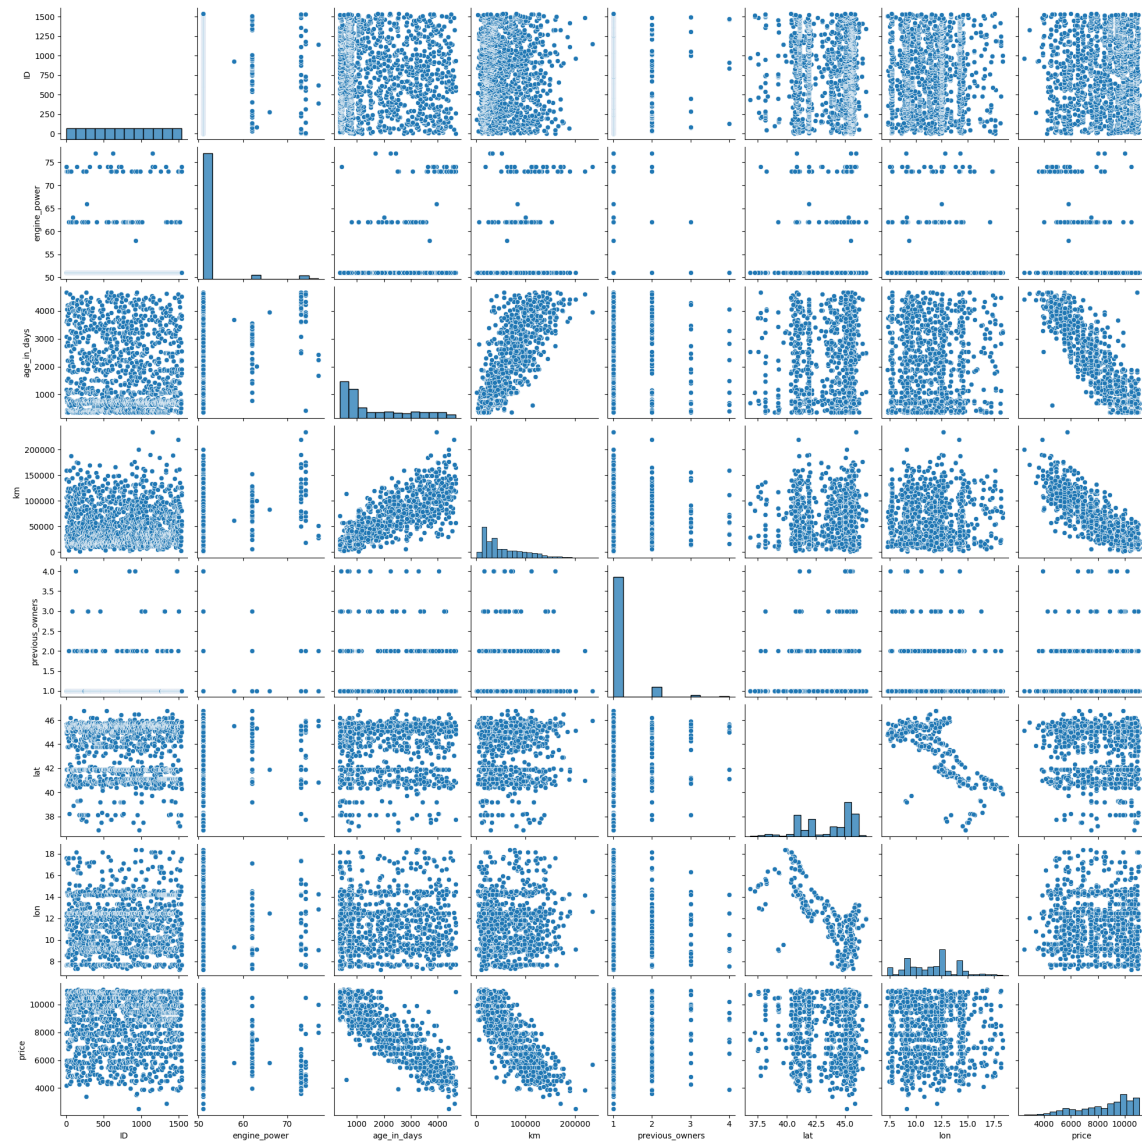
In [6]:

#EDA

sns.pairplot(df)

Out[6]:

<seaborn.axisgrid.PairGrid at 0x11d927903a0>




In [7]:

```
sns.displot(['Price'])
```



Out[7]:

```
<seaborn.axisgrid.FacetGrid at 0x11d97ed2f80>
```




In [8]:

```
sns.displot(df['price'])
```



Out[8]:

```
<seaborn.axisgrid.FacetGrid at 0x11d97ed33d0>
```




In [9]:

```
sns.displot(df['engine_power'])
```



Out[9]:

```
<seaborn.axisgrid.FacetGrid at 0x11d92790070>
```




In [10]:

```
fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
          'lat', 'lon', 'price']]  
sns.heatmap(fiatdf.corr())
```

Out[10]:

<Axes: >




In [11]:

```
fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
          'lat', 'lon']]  
sns.heatmap(fiatdf.corr())#without price
```

Out[11]:

<Axes: >



In [12]:

```
X=fiatdf[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
          'lat', 'lon']]  
y=df['price']
```

In [13]:

```
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=101)  
from sklearn.linear_model import LinearRegression  
regr=LinearRegression()  
regr.fit(X_train,y_train)  
print(regr.intercept_)
```

8971.195683500027

In [14]:

```
coeff_df=pd.DataFrame(regr.coef_,X.columns,columns=['coefficient'])  
coeff_df
```

Out[14]:


	coefficient
ID	-0.046704
engine_power	11.646408
age_in_days	-0.898018
km	-0.017232
previous_owners	26.400886
lat	32.189709
lon	0.161073

In [15]:

```
predictions=regr.predict(X_test)  
plt.scatter(y_test,predictions)
```

Out[15]:

<matplotlib.collections.PathCollection at 0x11d9b303e80>




In [17]:

```
sns.displot((y_test,predictions),bins=50)#without semicolon
```


Out[17]:

```
<seaborn.axisgrid.FacetGrid at 0x11d9b2d2c80>
```



In [18]:

```
sns.displot((y_test,predictions),bins=50);#with semicolon
```



In [19]:

```
from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 593.0876179519989

MSE: 551442.6799691883

MAE: 742.5918663500081

In [20]:

```
#accuracy
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.8597136704308846

In [22]:

```
df.fillna(method='ffill',inplace=True)
```

In [23]:

```
x=np.array(df['age_in_days']).reshape(-1,1)  
y=np.array(df['km']).reshape(-1,1)  
df.dropna(inplace=True)
```

In [24]:

```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)  
regr.fit(X_train,y_train)  
regr.fit(X_train,y_train)
```

Out[24]:


```
LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [25]:

```
y_pred=regr.predict(X_test)  
plt.scatter(X_test,y_test,color='y')  
plt.plot(X_test,y_pred,color='b')  
plt.show()
```



In [57]:



```
#elasticnet
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
[-90.71386288]
8677.92428806976
Mean Squared Error on test set 3678763.554360448
```

In []:

